

The Relational Data Model

2. The Relational Data Model

- use a simple and uniform data structure: the relation
- has been implemented in most commercial database systems
- has a solid theoretic foundation.

2.1 Structures

- In the relational model, everything is described using relations.
- A relation can be thought of as a named table.
- Each column of the table corresponds to a named attribute.
- The set of allowed values for an attribute is called its domain.
- Each row of the table is called a tuple of the relation.
- N.B. There is no ordering of column or rows.

Example

PLAYER					
Name	Position	Goals	Age	Height	Weight
Heady	Half-forward	17	24	183	83
Sumich	Full-forward	59	26	191	92
Langdon	Utility	23	23	189	86

PLAYER					
Name	Age	Height	Weight	Goals	Position
Sumich	26	191	92	59	Full-forward
Langdon	23	189	86	23	Utility
Heady	24	183	83	17	Half-forward

Above two tables are the same relation ---- Player

- Mathematically,
 - a *domain* D is a set of atomic values (having some fixed data type) which represent some semantic meaning.
 - an *attribute*, A , is the name of a role played by a *domain*, $dom(A)$.
 - a *relation schema* R , denoted by $R(A_1, A_2, \dots, A_n)$, is a set of attributes $R = \{A_1, A_2, \dots, A_n\}$.

Composite and multivalued attributes are disallowed!

- A *tuple*, $t(A_1, A_2, \dots, A_n)$, is a point in $dom(A_1) \times \dots \times dom(A_n)$ where each $dom(A_j)$ is the domain of A_j .
- A *relation* (or a *relation instance*) is a set of tuples: a subset of $dom(A_1) \times \dots \times dom(A_n)$.
- A relation schema is used to describe a relation.
- The *degree* of a relation is the number of attributes of its relation schema.

Relational Data Model vs ER Model:

- Relation schema (intension) \rightleftharpoons entity or relationship type schema (intension).
- attributes \rightleftharpoons attributes
- tuple \rightleftharpoons instance of entity/relationship
- relation (instance, extension) \rightleftharpoons entity/relationship extension
- composite and multivalued attributes are allowed in ER model, but not allowed in relational data model.

- *Keys* are used to identify tuples in a relation.
- A *superkey* is a set of attributes that uniquely determines a tuple.
- Note that this is a property of the relation that does not depend on the current relation instance.
- A candidate key is a superkey, none of whose *proper* subsets is a superkey.
- Keys are determined by the applications.
- E.g. if {Name} is unique then it is a candidate key for PLAYER; otherwise we need to use the whole tuple or create a candidate key, say PID.
- {Goals} usually cannot not be a candidate key since different players *might* have the same number of goals.
- {Name, Goals} is a superkey but not a candidate key if {Name} is a key.

- A *primary key* is a designated candidate key.
- In many applications it is necessary to invent a primary key if there is no natural one - often this would be a non-negative integer
- e.g. Person_number.
- When a relation schema has several candidate keys, usually better to choose a primary key with a single attribute or a small number of attributes.

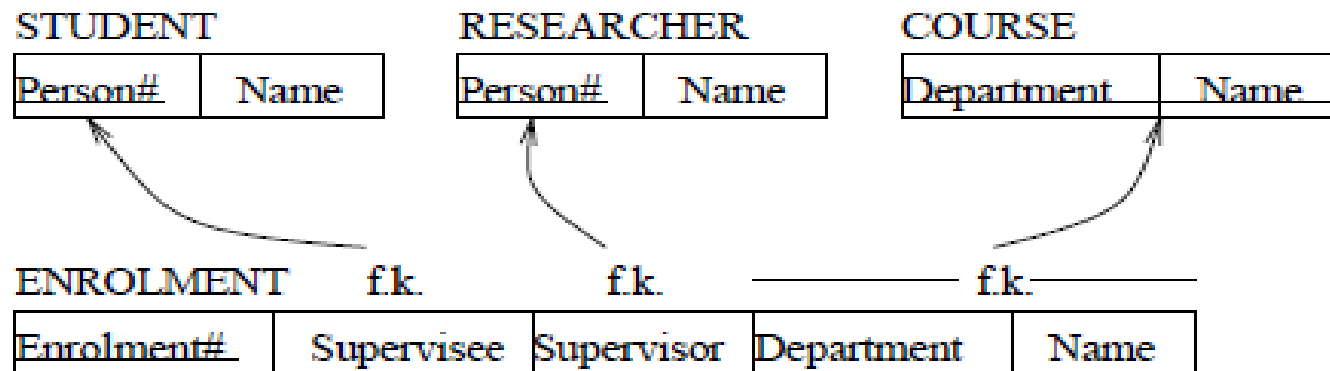
2.2 Integrity constraints

- There are several kinds of integrity constraints that are an integral part of the relational model:
- **2.2.1 Key constraint:** candidate key values must be unique for every relation instance.
- **2.2.2 Entity integrity:** an attribute that is part of a primary key cannot be NULL.
- **2.2.3 Referential integrity:** The third kind has to do with “foreign keys”.

- Foreign keys are used to refer to a tuple in another relation.
- A set, FK , of attributes from a relation schema R_1 may be a foreign *key* if
 - the attributes have the same domains as the attributes in the primary key of another relation schema R_2 , and
 - a value of FK in a tuple t_1 of R_1 either occurs as a value of PK for some tuple t_2 in R_2 or is null.
- *Referential integrity*: The value of FK must occur in the other relation or be entirely NULL.

2.2.4 Checking constraints on updates

- To maintain the integrity of the database, we need to check that integrity constraints will not be violated before proceeding with an update.
- Example: Suppose we have the following schema with foreign keys as shown:



<2, Dr. V. Ciesielski>

insert

RESEARCHER	
Person#	Name
1	Dr.C.C.Chen
2	Dr.R.G.Wilkinson

STUDENT	
Person#	Name
1	Dr.C.C.Chen
3	Ms.K.Juliff
4	Ms.J.Gledill
5	Ms.B.K.Lee

COURSE	
Department	Name
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT				
Enrolment#	Supervisee	Supervisor	Department	Name
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

<Comp.Sci., NULL>

insert

STUDENT	
Person#	Name
1	Dr.C.C.Chen
3	Ms.K.Juliff
4	Ms.J.Gledill
5	Ms.B.K.Lee

RESEARCHER	
Person#	Name
1	Dr.C.C.Chen
2	Dr.R.G.Wilkinson

COURSE	
Department	Name
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT				
Enrolment#	Supervisee	Supervisor	Department	Name
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

<5, 6, 2, Psychology, Ph.D>

insert

STUDENT	
Person#	Name
1	Dr.C.C.Chen
3	Ms.K.Juliff
4	Ms.J.Gledill
5	Ms.B.K.Lee

RESEARCHER	
Person#	Name
1	Dr.C.C.Chen
2	Dr.R.G.Wilkinson

COURSE	
Department	Name
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT				
Enrolment#	Supervisee	Supervisor	Department	Name
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

STUDENT	
Person#	Name
1	Dr.C.C.Chen
3	Ms.K.Juliff
4	Ms.J.Gledill
5	Ms.B.K.Lee

RESEARCHER	
Person#	Name
1	Dr.C.C.Chen
2	Dr.R.G.Wilkinson

COURSE	
Department	Name
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT				
Enrolment#	Supervisee	Supervisor	Department	Name
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

- *Insertions*: When inserting, we need to check
 - that the candidate keys are not already present,
 - that the value of each foreign key either
 - is all null, or
 - is all non-NULL and occurs in the referenced relation.

Examples:

1. Insert $\langle 2, Dr.V.Ciesielski \rangle$ into RESEARCHER

Allowed? No. Violates a key constraint.

Action? Reject or allow the user to correct.

2. Insert $\langle \textit{Comp.Sci.}, \textit{NULL} \rangle$ into COURSE

Allowed? No. Violates the entity integrity constraint.

Action: Reject or correct.

3. Insert $\langle 5, 6, 2, \textit{Psychology}, \textit{Ph.D.} \rangle$ into
ENROLMENT

Allowed? No. Violates a referential integrity constraint (There is no person number 6).

Action: Reject, correct or accept after insertion of person number 6.

- *Deletions*: When deleting, we need to check referential integrity – check whether the primary key occurs in another relation.

Examples:

1. Delete tuple with Person# = 2 from RESEARCHER

Allowed? No. Violates the referential integrity.

Action: Reject, correct or modify the ENROLMENT tuple by

- deleting it (note that this requires another integrity check, possibly causing a cascade of deletions), or
- setting the foreign key value to NULL (note this can't be done if it is part of a primary key), or
- setting the foreign key value to another acceptable value.

Modifications:

If the modified attribute is a

- primary key: this is similar to deleting and then reinserting.
- foreign key: check that the new value refers to an existing tuple.
- neither: no problems can arise.

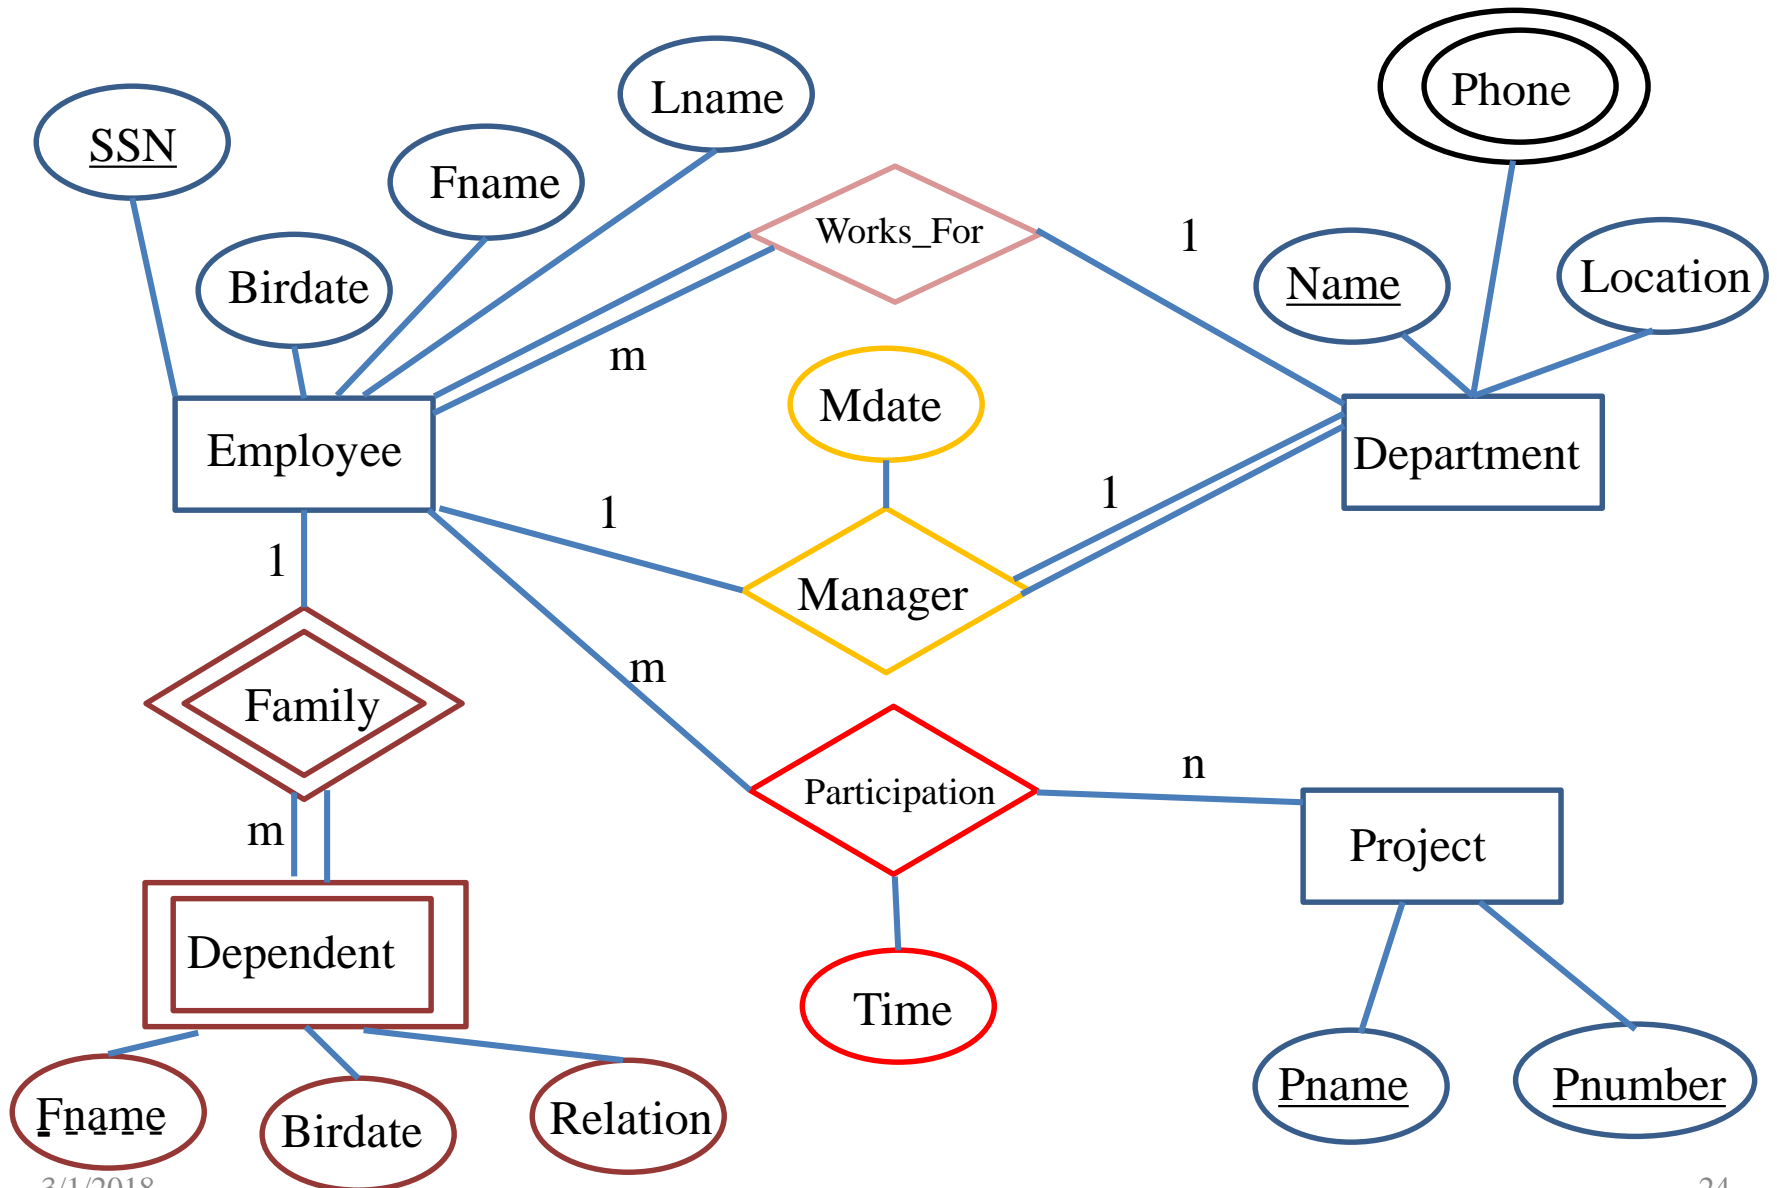
2.2.5 Relational database definition

- A *relational database schema*, is a set of relation schema $\{R_1, \dots, R_m\}$ and a set of integrity constraints.
- A relational database instance is a set of relation instances $\{r_1, \dots, r_m\}$ such that each r_i is an instance of R_i , and the integrity constraints are satisfied.

2.3 ER to Relational Data Model Mapping

- One technique for database design is to first design a conceptual schema using a high-level data model, and then map it to a conceptual schema in the DBMS data model for the chosen DBMS.
- Here we look at a way to do this mapping from the ER to the relational data model.
- It involves the following 7 steps.

- Example: ER \rightarrow RDB



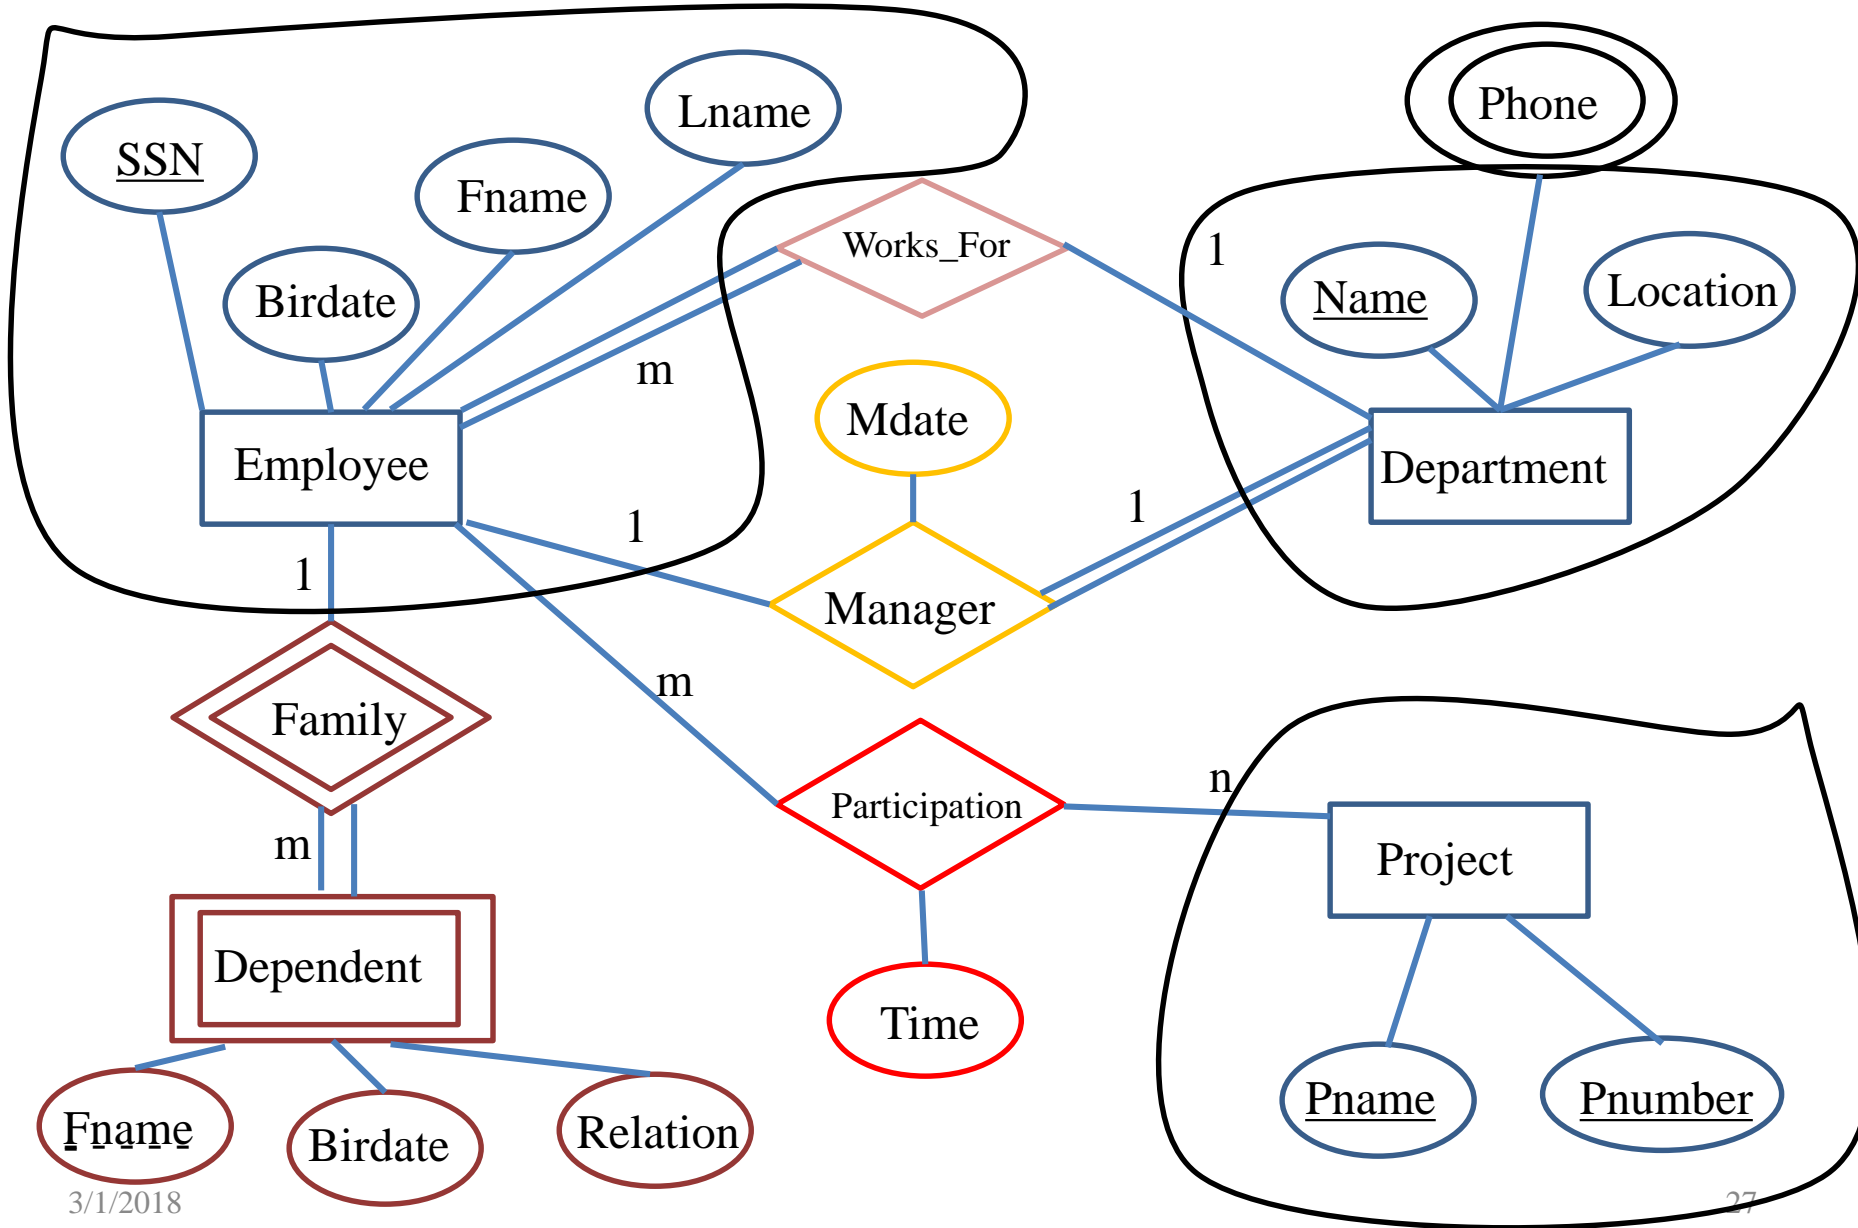
*

*

- Step 1 : For each regular (not weak) entity type E, create a relation R with
 - Attributes : All simple attributes (and simple components of composite attributes) of E.
 - Key : Choose one of the keys of E as the primary key for the relation.

- Step 1a : For each specialised entity type E, with parent entity type P, create a relation R with
 - Attributes : The attributes of the key of P, plus the simple attributes of E.
 - Key : The key of P.

- Example: ER \rightarrow RDB



Employee

<u>SSN</u>	Fname	Lname	Birdate
------------	-------	-------	---------

Department

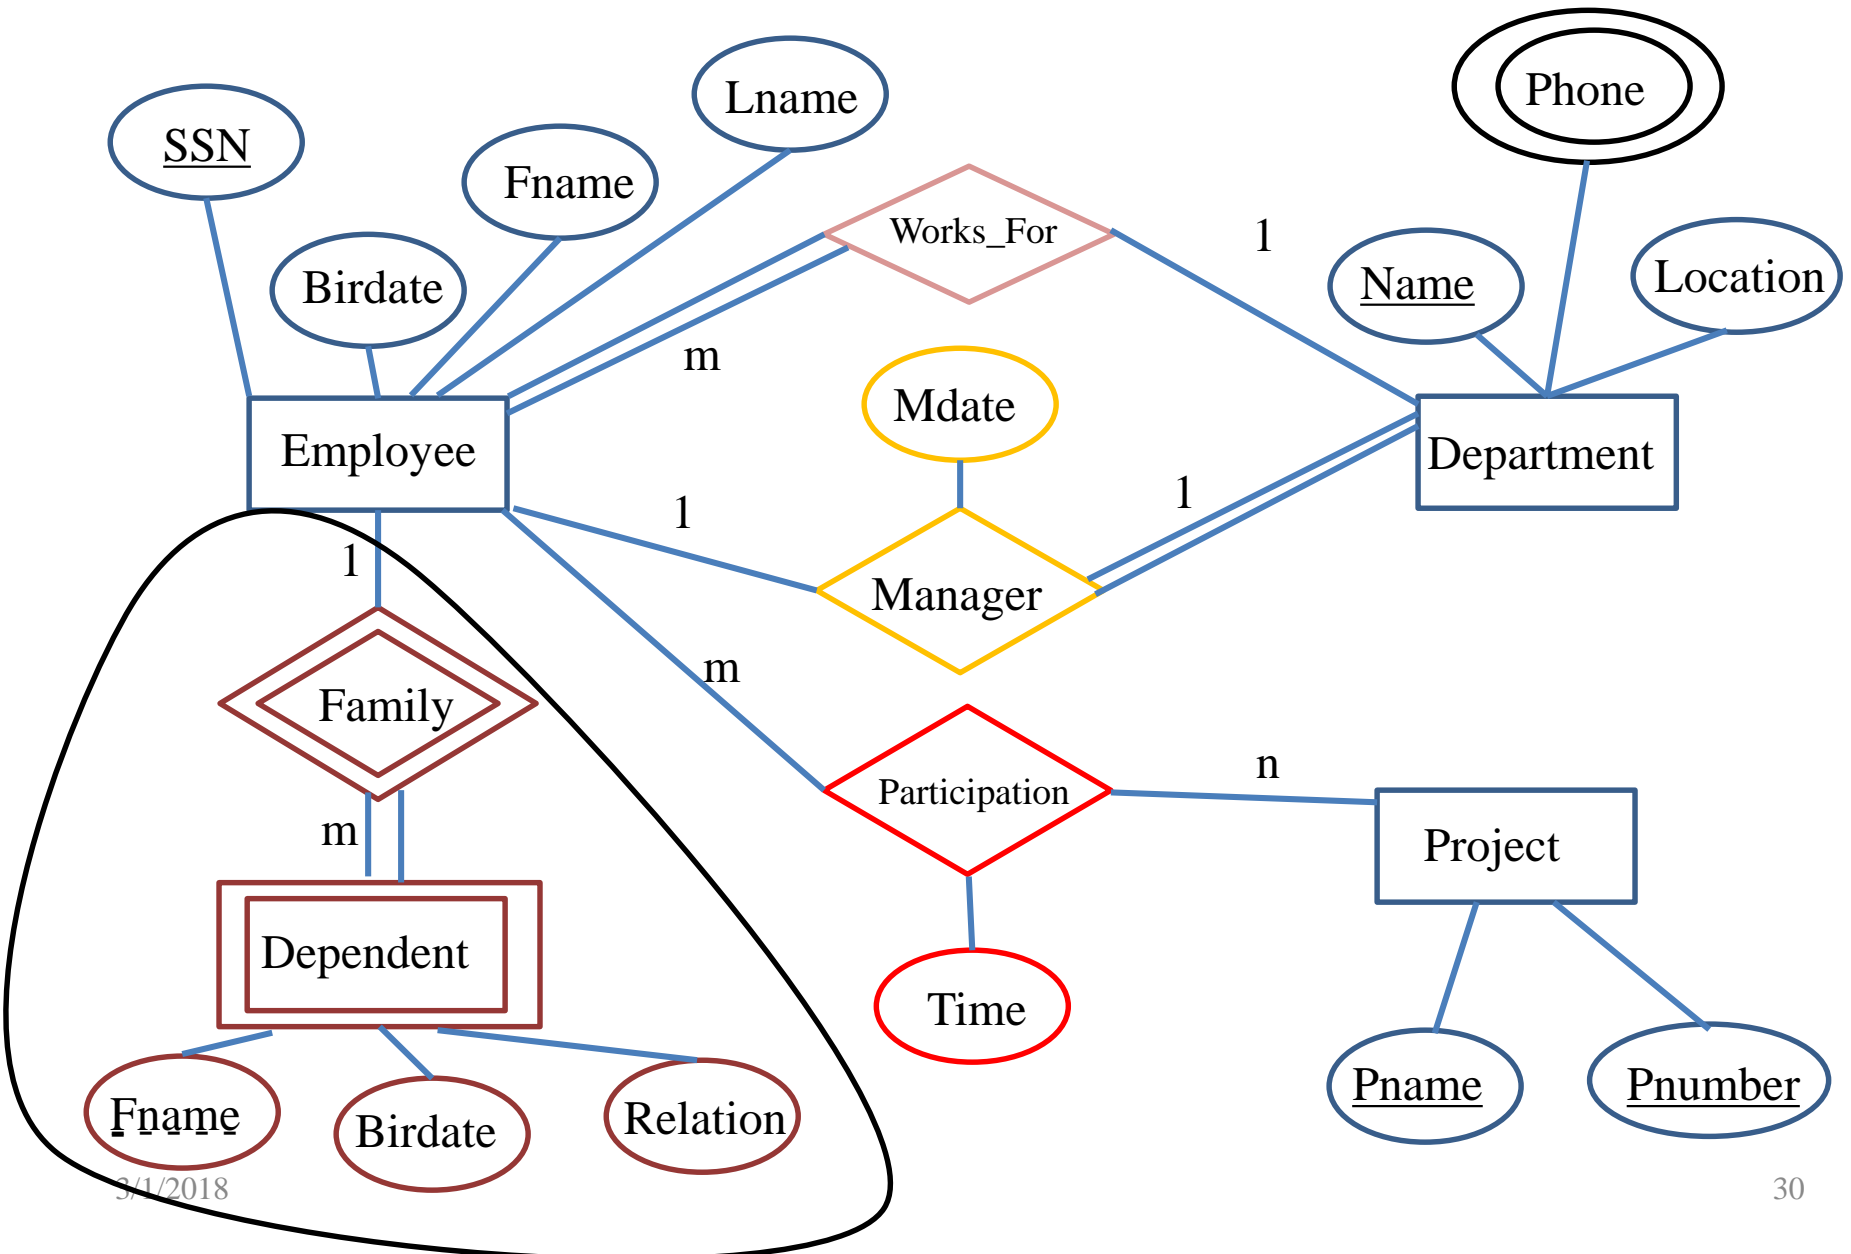
<u>Name</u>	Location
-------------	----------

Project

<u>Pname</u>	<u>Pnumber</u>
--------------	----------------

- Step 2 : For each weak entity type W, with owner entity type E, create a relation R with
 - Attributes : All simple attributes (and simple components of composite attributes) of W, and include as a foreign key the prime attributes of the relation derived from E.
 - Key : The foreign key plus the partial key of W.

- Example: ER \rightarrow RDB



Employee

<u>SSN</u>	Fname	Lname	Birdate
------------	-------	-------	---------

Department

<u>Name</u>	Location
-------------	----------

Project

<u>Pname</u>	<u>Pnumber</u>
--------------	----------------

Dependent

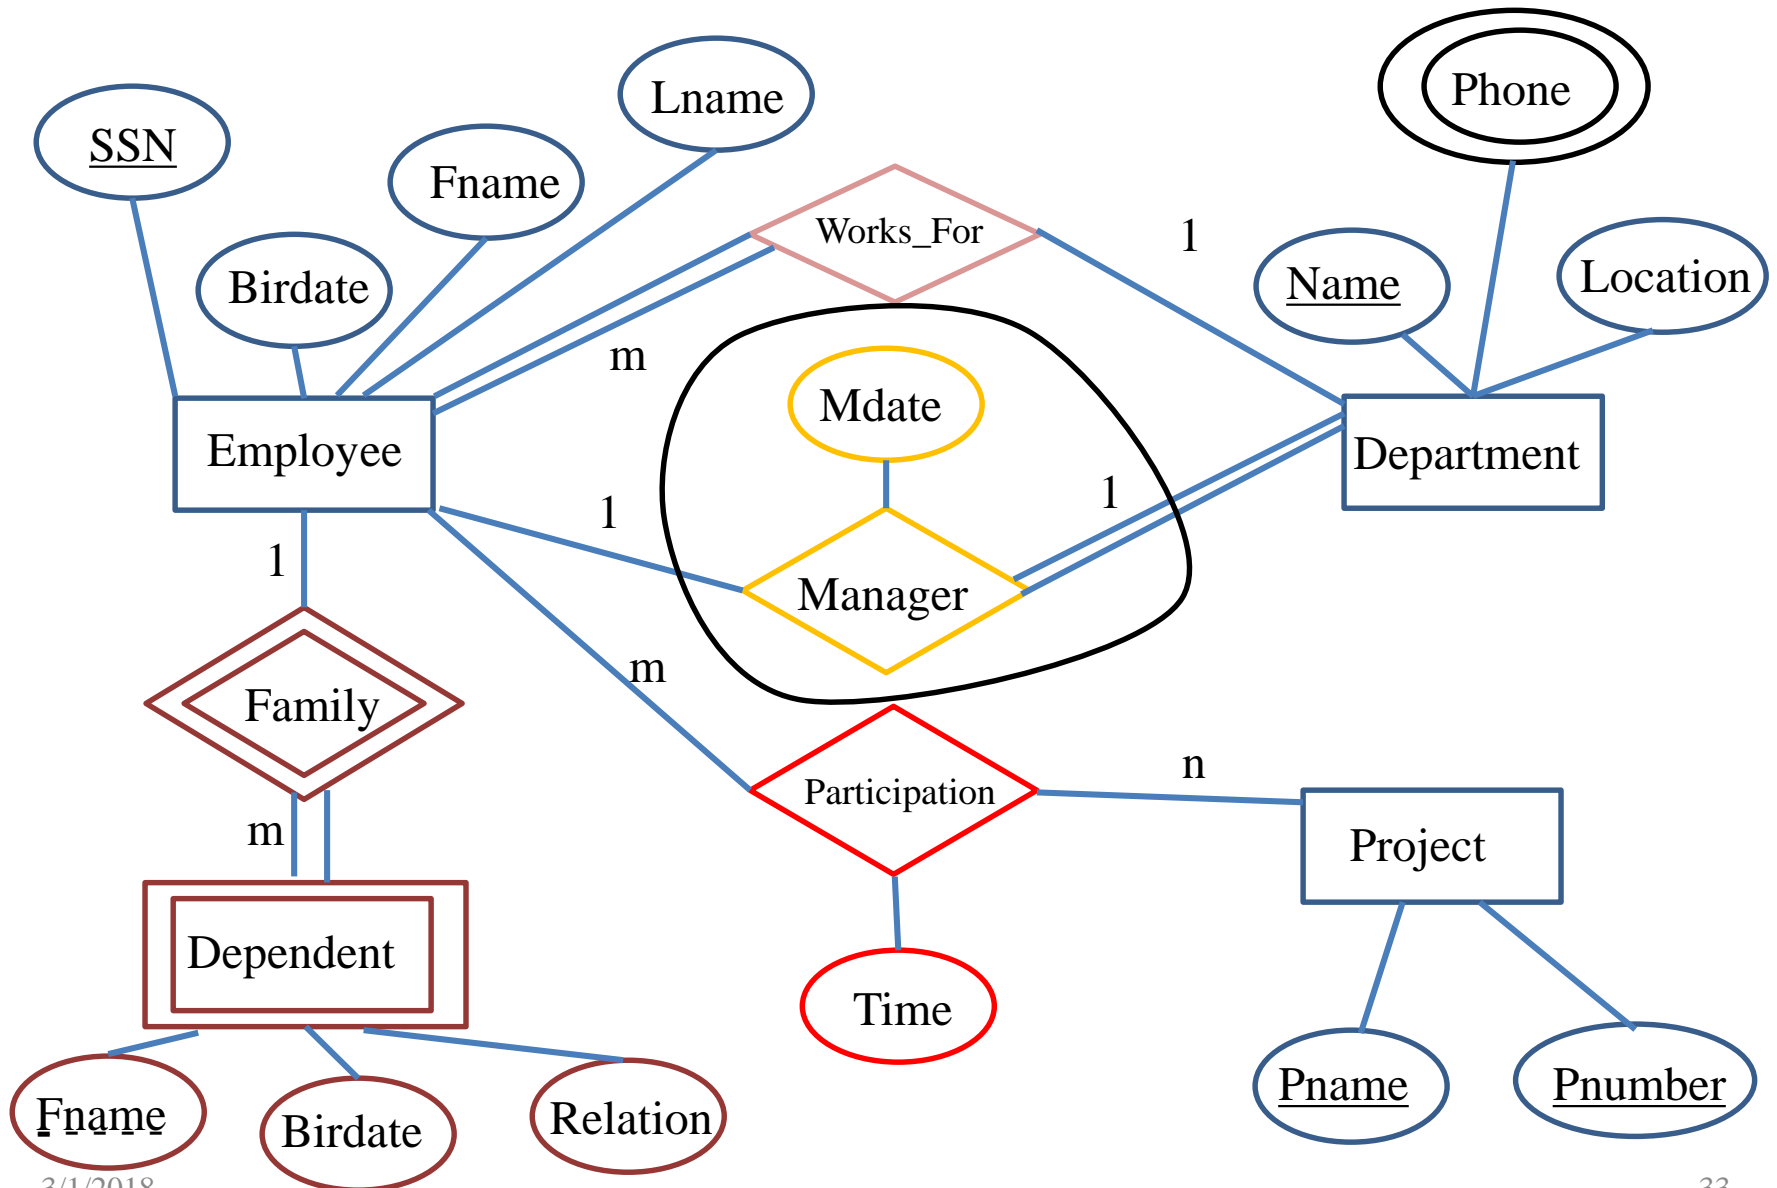
<u>SSN</u>	<u>Fname</u>	Birdate	Relation
------------	--------------	---------	----------



- Step 3 : For each 1:1 relationship type B. Let E and F be the participating entity types. Let S and T be the corresponding relations.
 - Choose one of S and T (prefer one that participates totally), say S.
 - Add the attributes of the primary key of T to S as a foreign key.
 - Add the simple attributes (and simple components of composite attributes) of B as attributes of S.

(Alternative: merge the two entity types and the relationship into a single relation, especially if both participate totally and do not participate in other relationships).

- Example: ER \rightarrow RDB



Employee

<u>SSN</u>	Fname	Lname	Birdate
------------	-------	-------	---------

Department

<u>Name</u>	Location	MSSN	Mdate
-------------	----------	------	-------

Project

<u>Pname</u>	<u>Pnumber</u>
--------------	----------------

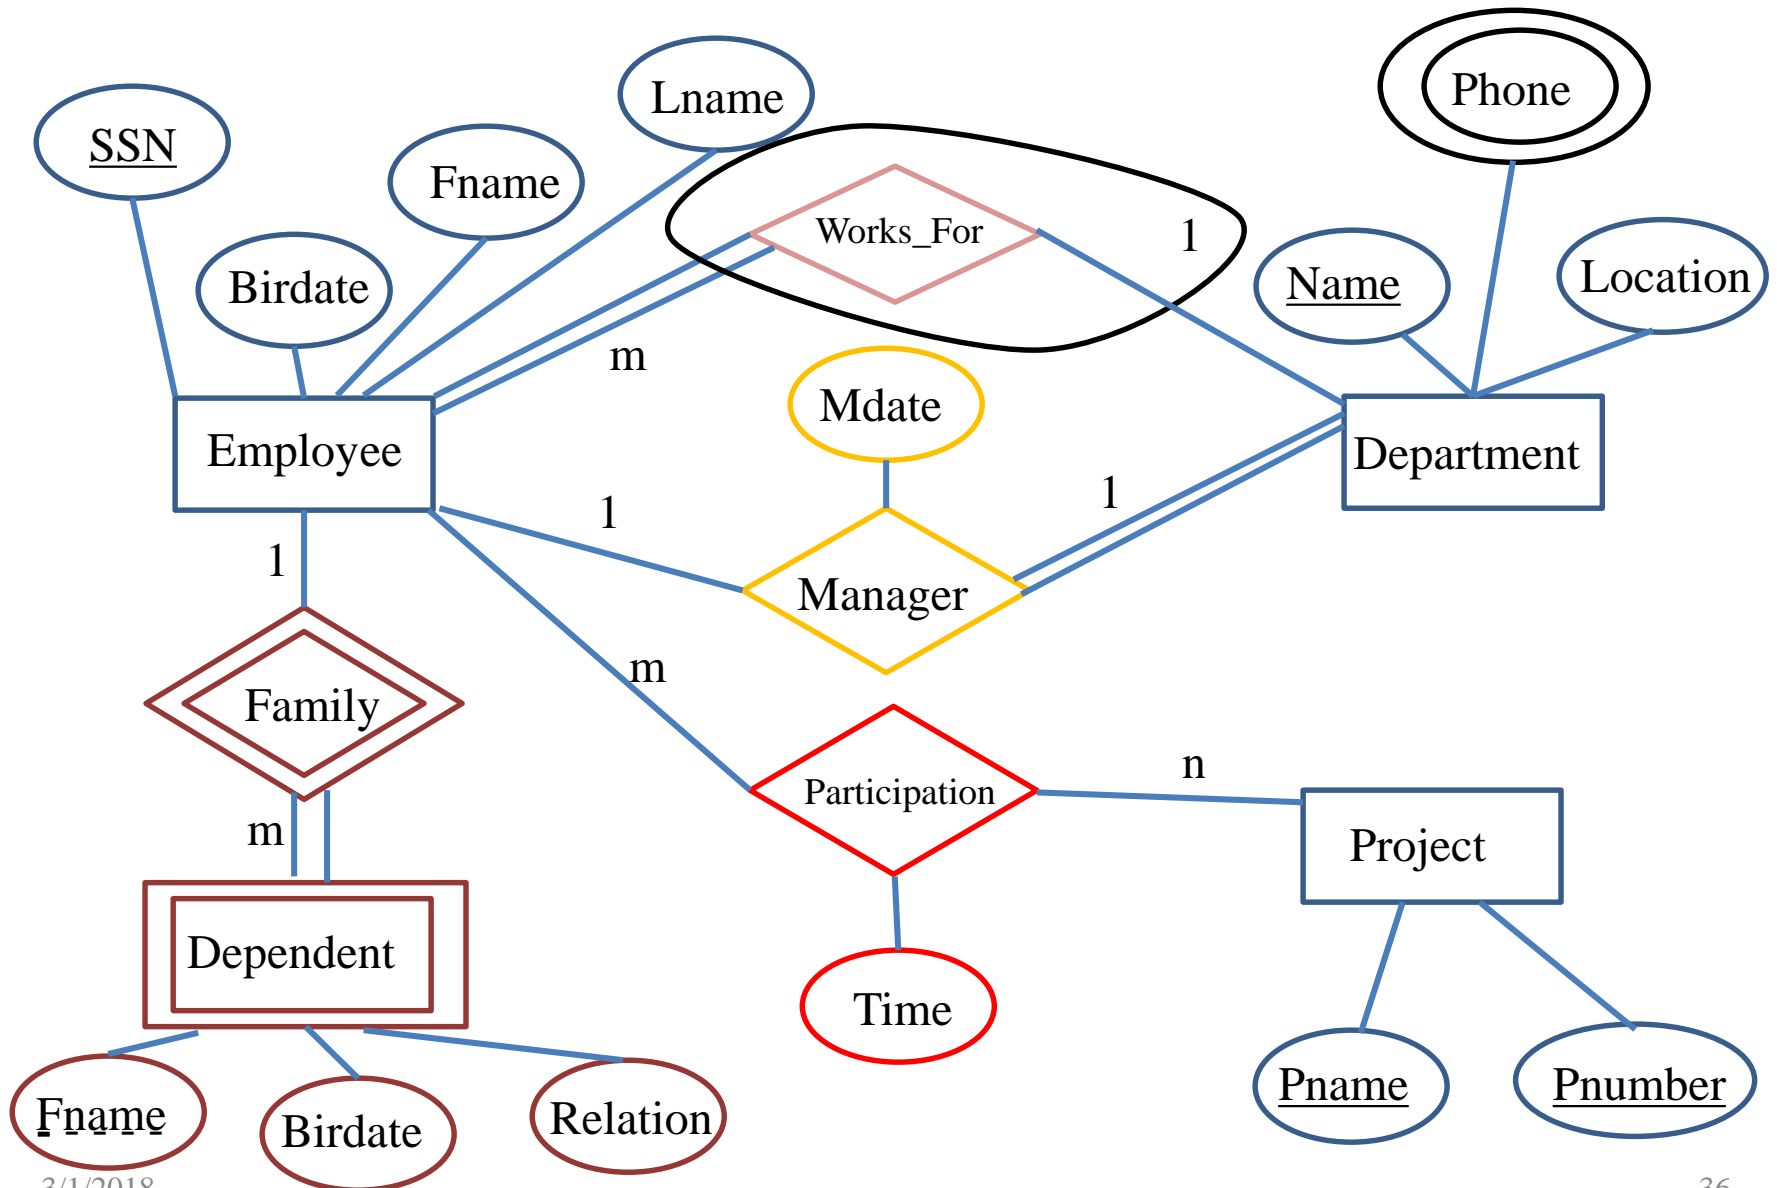
Dependent

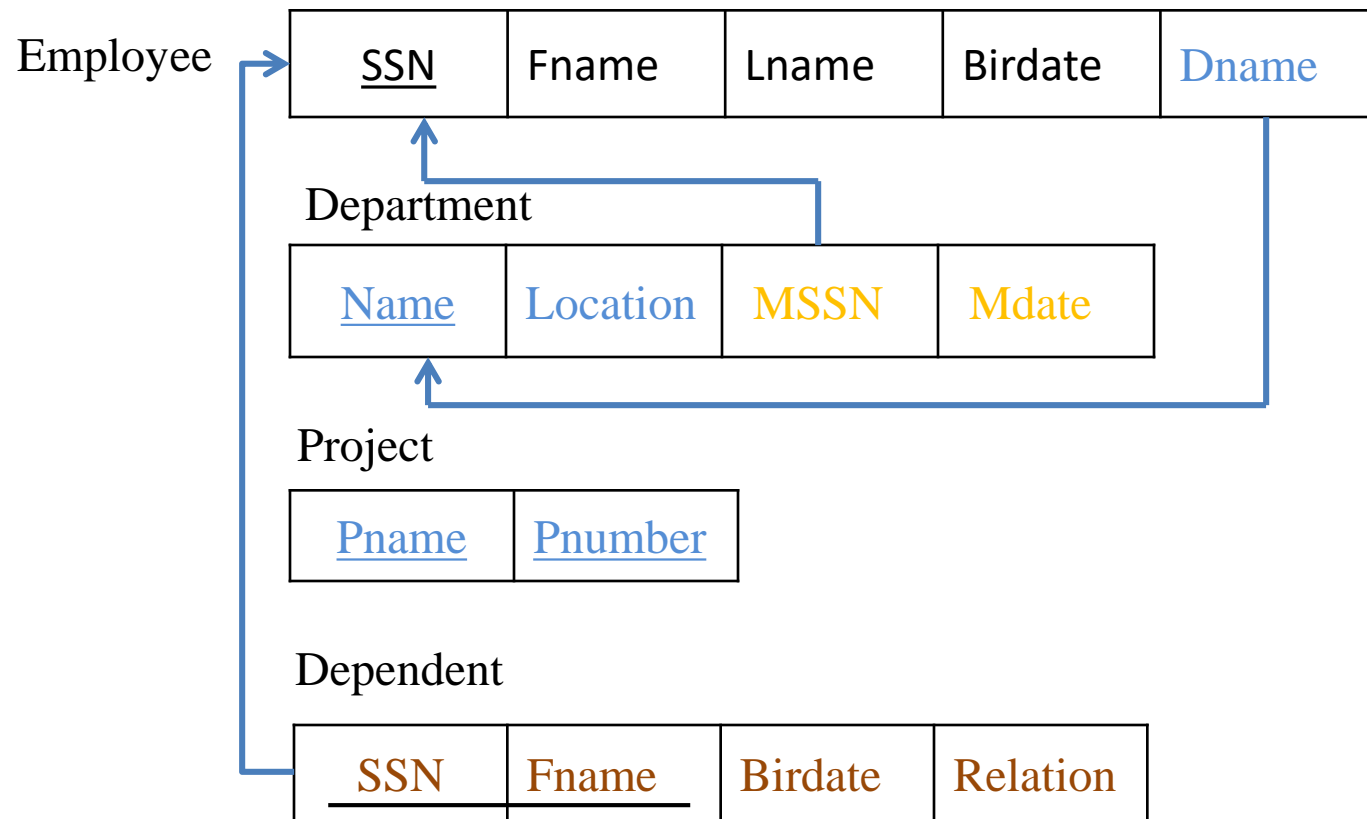
<u>SSN</u>	<u>Fname</u>	Birdate	Relation
------------	--------------	---------	----------

- Step 4 : For each regular 1:N relationship type B.
 - Let E and F be the participating entity types.
 - Let E be the entity type on the 1 side, F the one on the N side.
 - Let S and T be the corresponding relations.
 - Add the attributes of the primary key of S to T as a foreign key.
 - Add to T any simple attributes (or simple components of composite attributes) of the relationship.

(Notice that this doesn't add any new tuples, just attributes.)

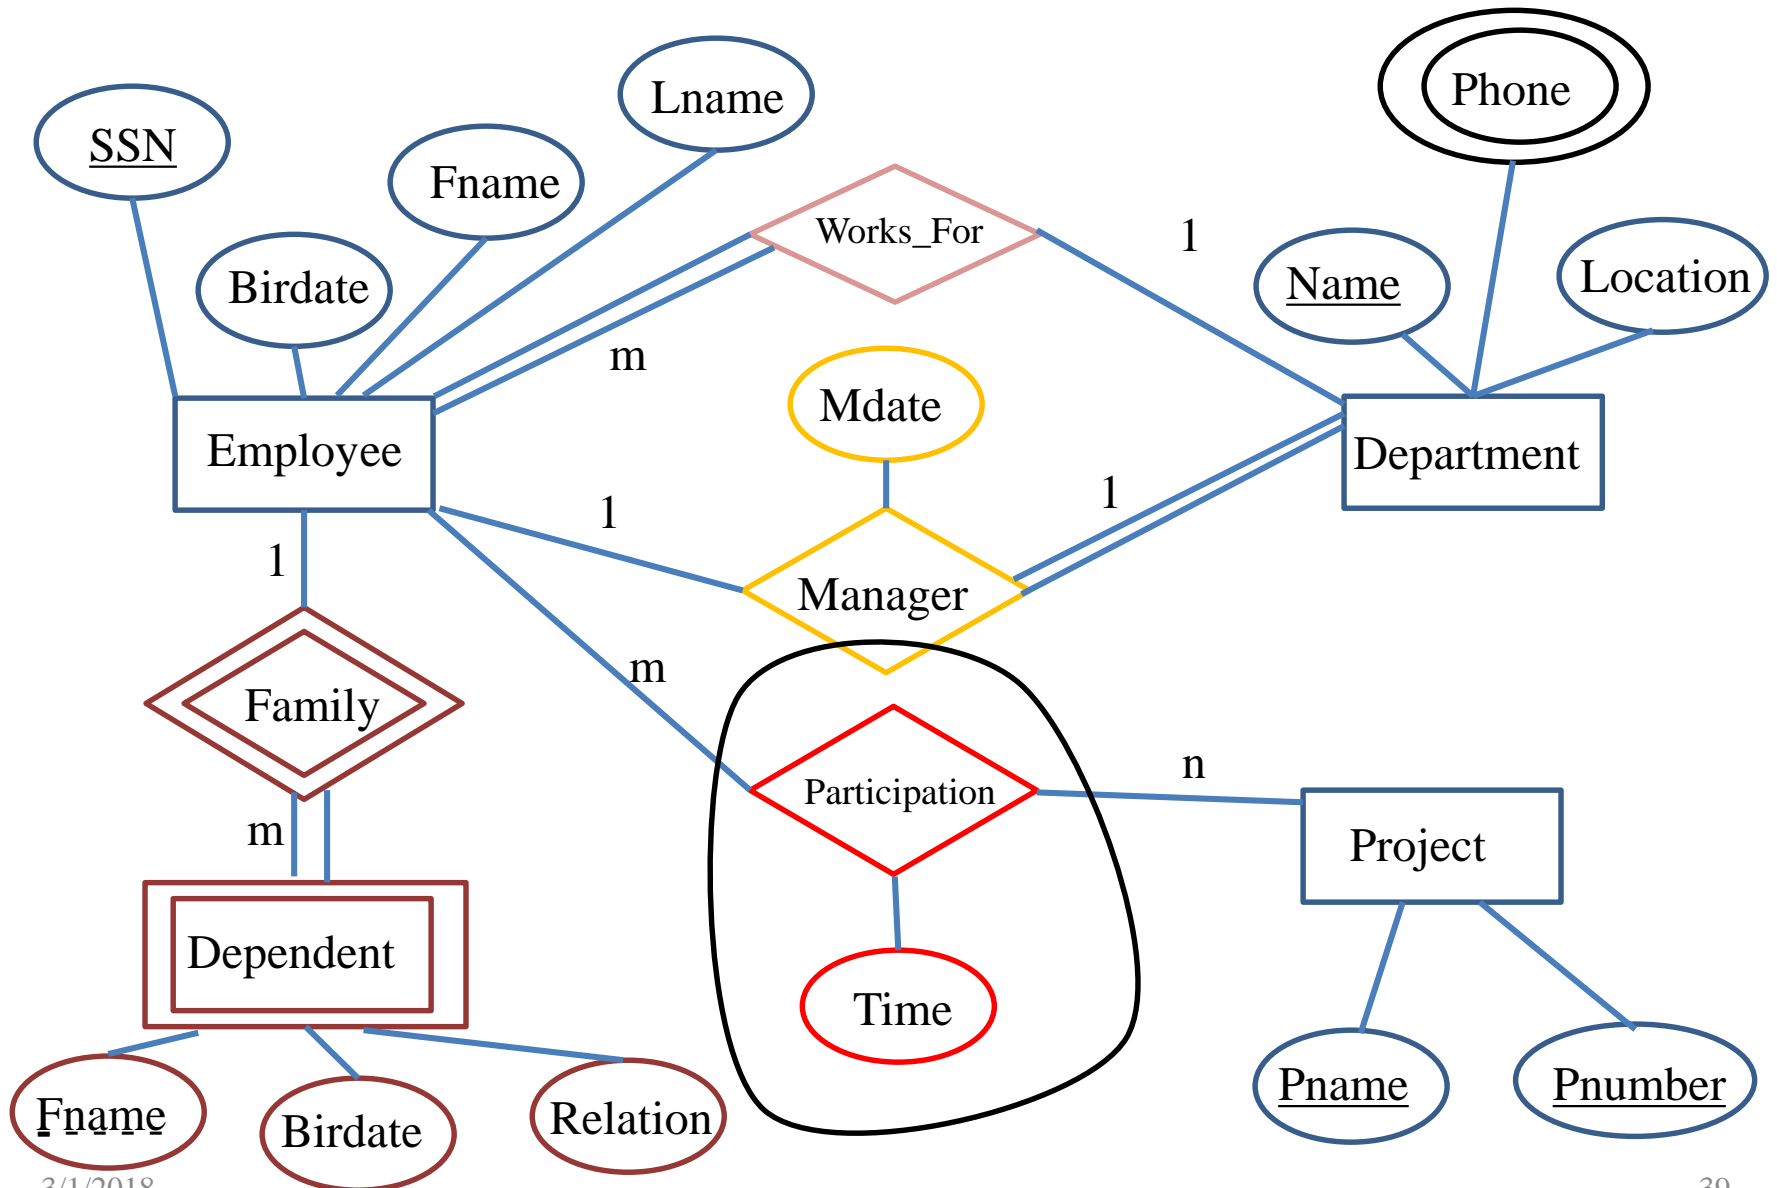
- Example: ER \rightarrow RDB

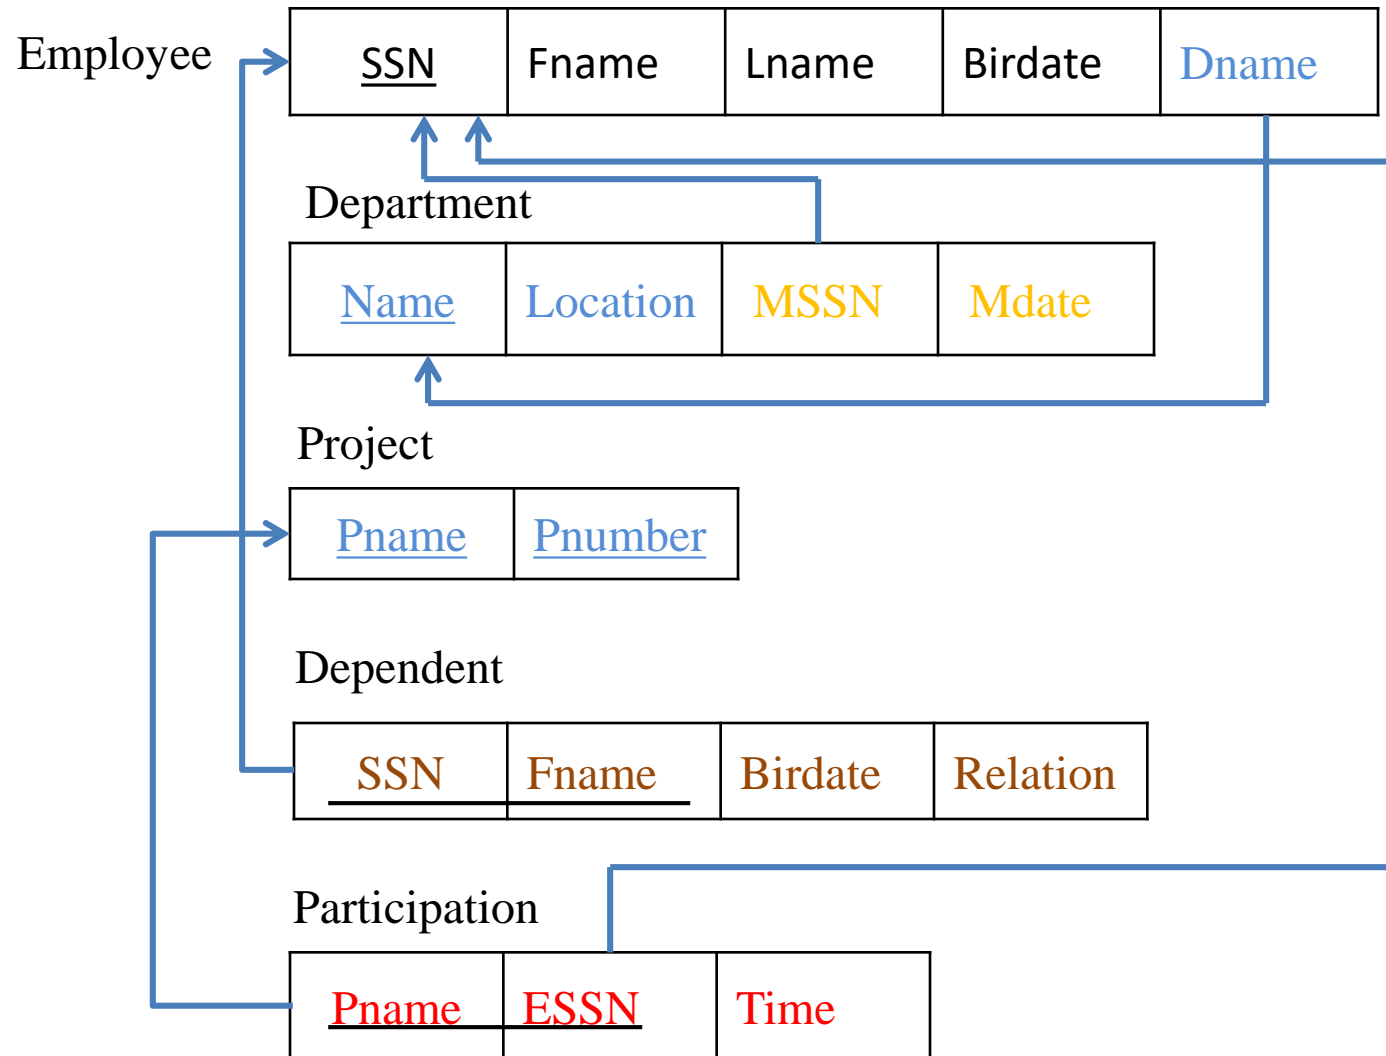




- Step 5 : For each N:M relationship type B. Create a new relation R. Let E and F be the participating entity types. Let S and T be the corresponding relations.
 - Attributes : The key of S and the key of T as foreign keys, plus the simple attributes (and simple components of composite attributes) of B.
 - Key : The key of S and the key of T.

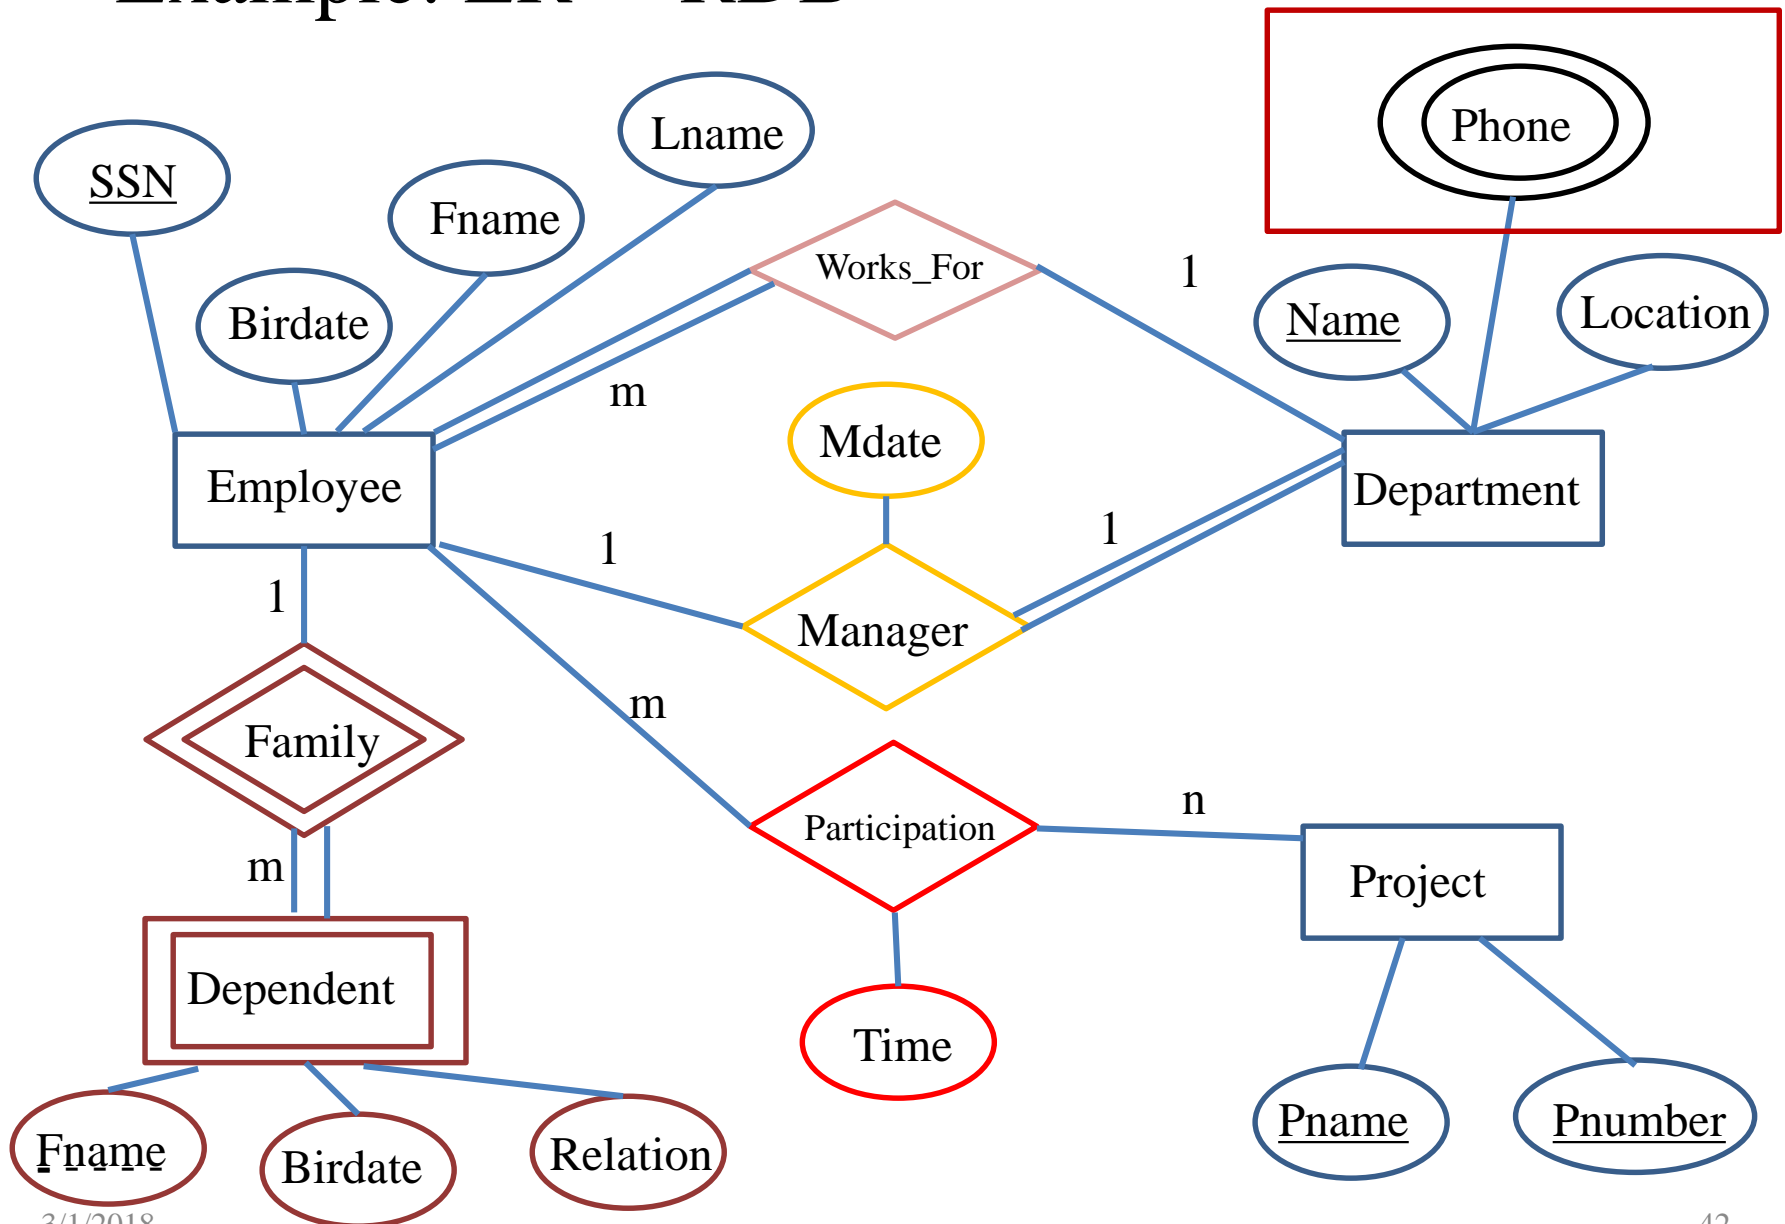
- Example: ER \rightarrow RDB



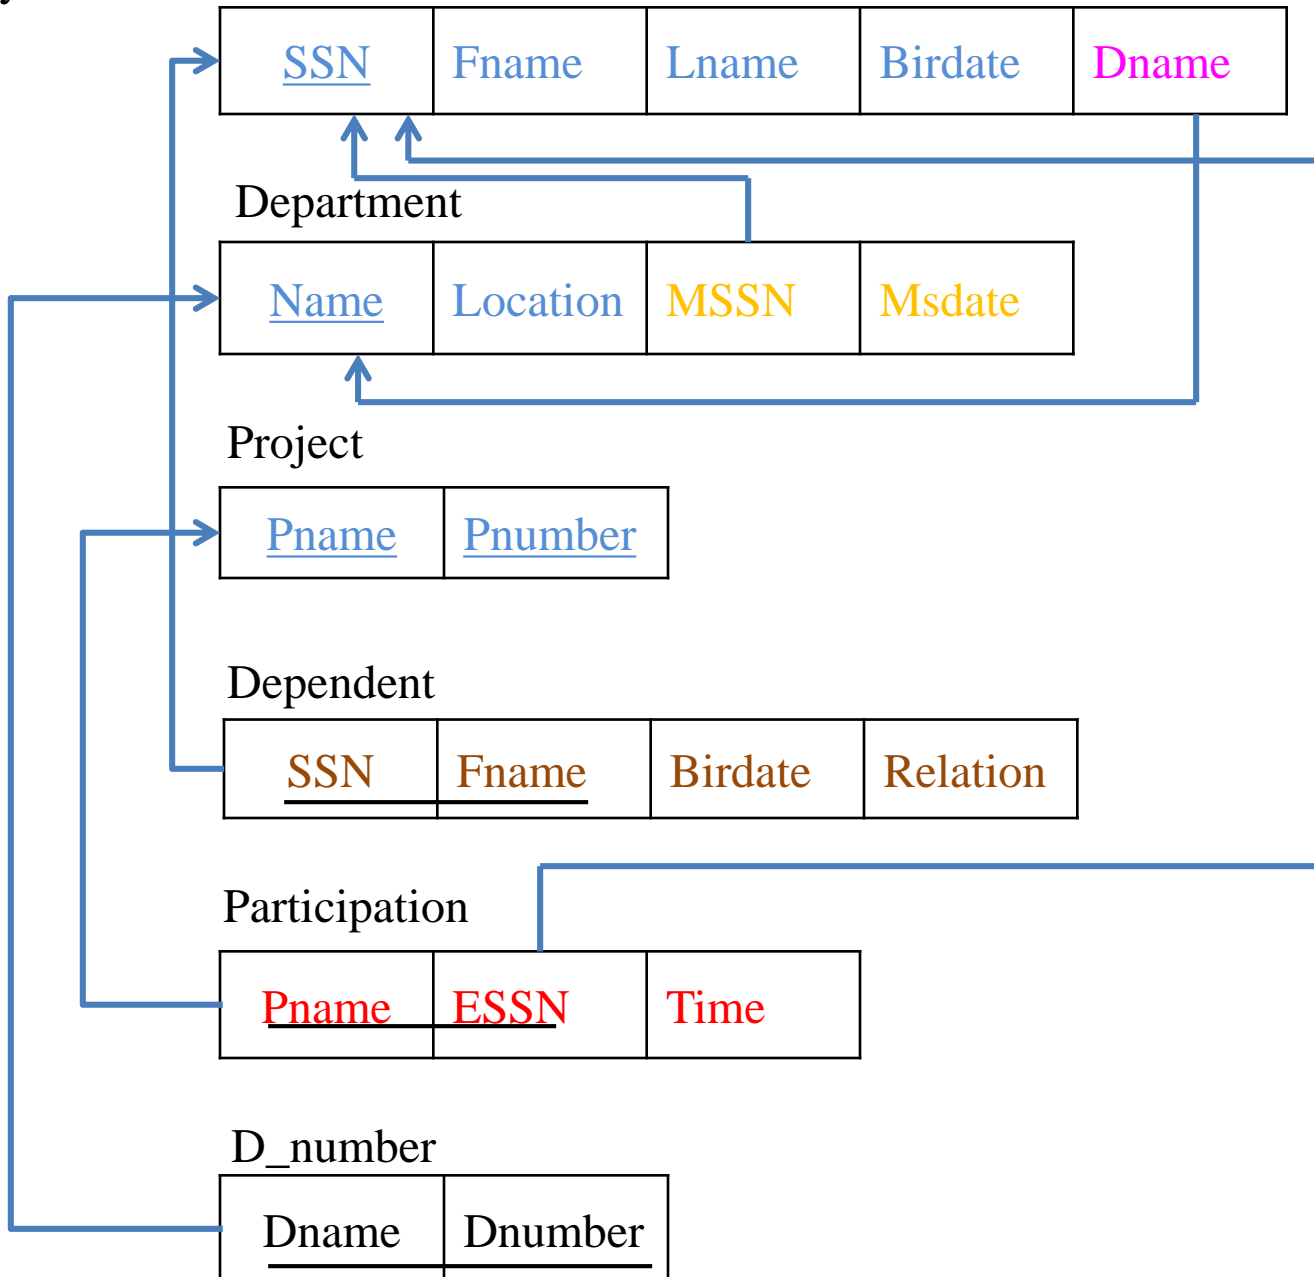


- Step 6 : For each multivalued attribute A. Create a new relation R. Let A be an attribute of E.
 - Attributes :
 1. A (if A is a simple attribute) together with the key of E as a foreign key.
 2. The simple components of A (if A is a composite attribute), together with the key of E as a foreign key.
 - Key : All attributes.

- Example: ER \rightarrow RDB



Employee



- Step 7 : For each n-ary relationship type ($n > 2$). Create a new relation with
 - Attributes : as for Step 5.
 - Key : as for Step 5, except that if one of the participating entity types has participation ratio 1, its key can be used as a key for the new relation.

